



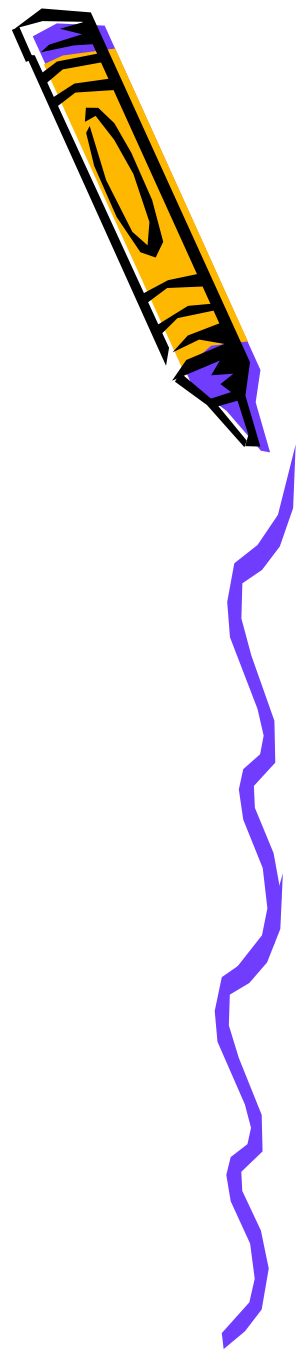
Computer Architecture

Hossein Asadi
Department of Computer Engineering
Sharif University of Technology
asadi@sharif.edu



Today's Topics

- Performance Evaluation
- Standard Benchmarks



Copyright Notice



- Parts (text & figures) of this lecture adopted from:
 - Computer Organization & Design, The Hardware/Software Interface, 4th Edition, by D. Patterson and J. Hennessey, MK publishing, 2012.
 - "Intro to Computer Architecture" handouts, by Prof. Hoe, CMU, Spring 2009.
 - "Computer Architecture & Engineering" handouts, by Prof. Kubiawicz, UC Berkeley, Spring 2004.
 - "Intro to Computer Architecture" handouts, by Prof. Hoe, UWisc, Spring 2019.
 - "Computer Arch I" handouts, by Prof. Garzarán, UIUC, Spring 2009.



Performance Metrics



- Latency
 - Time between start and finish of a single task
- Throughput
 - Number of tasks finished in a given unit of time
- Question:
 - $\text{Throughput} = 1 / \text{Latency} (?)$
 - $\text{Latency} = 1 / \text{Throughput} (?)$



Throughput vs. Latency



- Example: Boeing 747 vs. BAC Concorde
 - $\text{Speed}_{\text{Concorde}} > \text{Speed}_{\text{Boeing}}$
 - $\text{Latency}_{\text{Concorde}} < \text{Latency}_{\text{Boeing}}$
 - Throughput???
 - How to define throughput?

| Airplane | Passenger Capacity | Speed (mph) |
|------------|--------------------|-------------|
| Boeing 747 | 470 | 610 |
| Concorde | 132 | 1350 |



Throughput vs. Latency (cont.)



- Throughput = Passenger Capacity * Speed

| Airplane | Passenger Capacity | Speed (mph) | Passenger Throughput |
|------------|--------------------|-------------|----------------------|
| Boeing 747 | 470 | 610 | 286,700 |
| Concorde | 132 | 1350 | 178,200 |



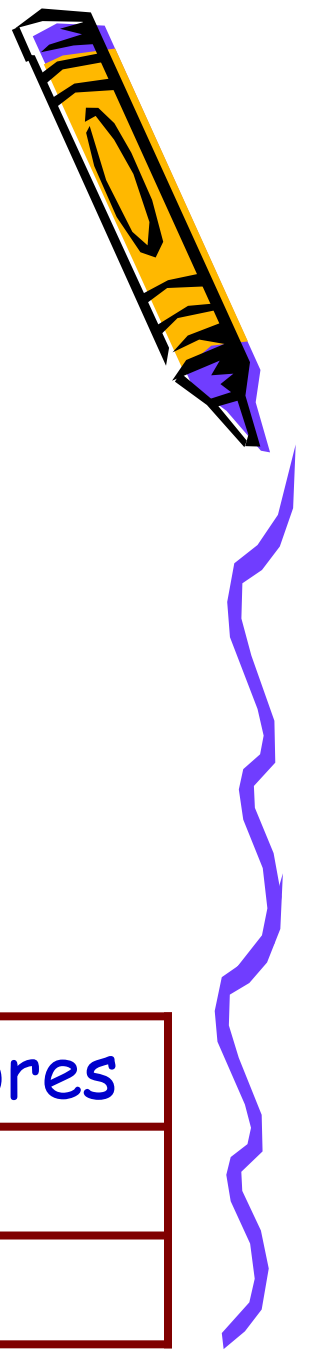
Throughput vs. Latency: Single-Core Example



- Consider Two Single-Core Computers
 - Less latency (response time)
 - The one that finishes tasks faster
 - More throughput
 - The one that finishes more jobs
 - Or the one that finishes tasks faster
 - In this example
 - Latency $\downarrow \Rightarrow$ Throughput \uparrow
 - Latency $\uparrow \Rightarrow$ Throughput \downarrow
 - Not always true!



Throughput vs. Latency: Multi-Core Example



- Computer A
 - Two cores, running at 3Mhz
- Computer B
 - Four cores, running at 2.5Mhz
- $\text{Latency}_A < \text{Latency}_B$
- $\text{Throughput}_A < \text{Throughput}_B$

| Computer | Frequency | # of cores |
|----------|-----------|------------|
| Type A | 3 Mhz | 2 |
| Type B | 2.5 Mhz | 4 |



Performance Definition



- Response Time
 - Total time to complete a task (program)
 - Also called, wall-clock time, elapsed time
- Performance = 1 / response time
 - Response time $\downarrow \Rightarrow$ Performance \uparrow
 - $\text{Performance}(x) / \text{Performance}(y) =$
 $\text{Execution time}(y) / \text{Execution time}(x)$



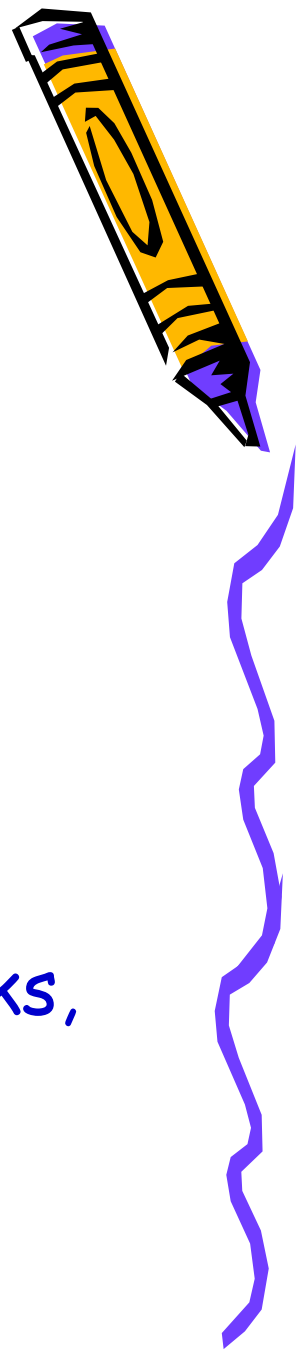
Performance Definition (cont.)



- Response Time Consists of:
 - CPU time
 - CPU time spent on a program
 - I/O time
 - Time elapsed to wait for I/O transactions
- CPU Time
 - User CPU time
 - CPU time directly spent on a program
 - System CPU time
 - CPU time spent in OS doing tasks on behalf of a program



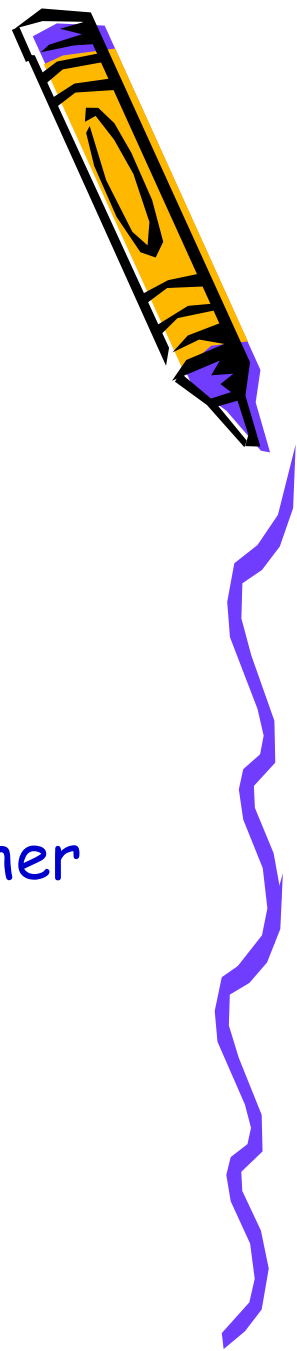
Performance Definition (cont.)



- System Performance
 - 1 / Elapsed time
- CPU Performance
 - 1 / User CPU time
- Clock Cycles
 - Clock periods, cycles, clock ticks, ticks, clocks



Response Time vs. Throughput



- Which One is More Important?
 - Response time or throughput
- User Perspective
 - Response time more visible
 - Unless a user runs bunch of tasks together
- System Admin Perspective
 - Throughput more visible



CPU Time

= CPU Clock Cycles * Clock Cycle Time

= CPU Clock Cycles / Clock Rate

- Example

- $\text{CPU}_A(\text{clock rate}) = 4 \text{ Ghz}$

- $\text{CPU}_B(\text{clock rate}) = 3 \text{ Ghz}$

- Which one runs program X faster?

- Depends on number of clock cycles spend on program X



CPU Time (cont.)

= CPU Clock Cycles * Clock Cycle Time

= CPU Clock Cycles / Clock Rate

- CPU Clock Cycles

 - = # of Instructions of a program * CPI

- CPI?

 - Average Clock Cycles per Instruction



CPU Time (cont.)

$$\begin{aligned} &= \text{Instr. Count} * \text{CPI} * \text{Clock Cycle Time} \\ &= (\text{Instr. Count} * \text{CPI}) / \text{Clock Rate} \end{aligned}$$

$$\text{Processor Performance} = \frac{\text{Time}}{\text{Program}}$$

$$\begin{array}{ccccc} \boxed{\begin{array}{c} \text{Instructions} \\ = \frac{\text{-----}}{\text{Program}} \end{array}} & \times & \boxed{\begin{array}{c} \text{Cycles} \\ \frac{\text{-----}}{\text{Instruction}} \end{array}} & \times & \boxed{\begin{array}{c} \text{Time} \\ \frac{\text{-----}}{\text{Cycle}} \end{array}} \\ \text{(code size)} & & \text{(CPI)} & & \text{(cycle time)} \end{array}$$

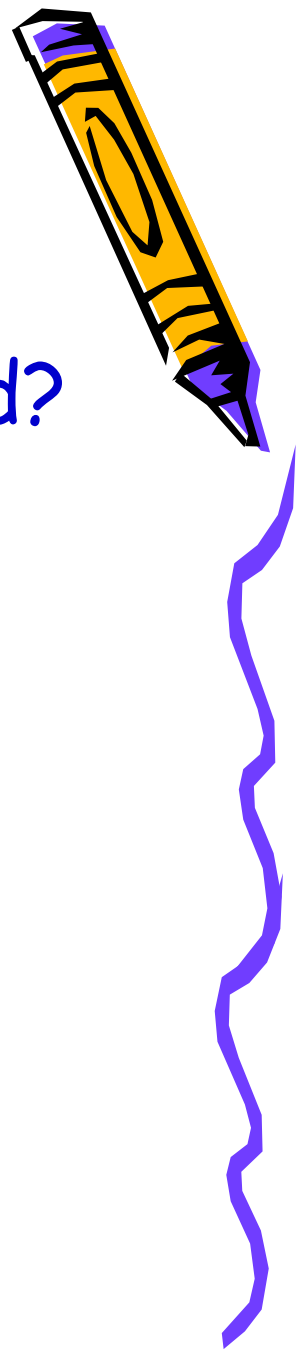
CPU Time (cont.)



- Example
 - Instruction count (IC) = 10,000
 - CPI = 4
 - Clock cycle time = 500ps or 0.5 ns
 - CPU time = $10,000 * 4 * 0.5 = 20 \text{ us}$
- Question:
 - IC same as code size or lines of code?



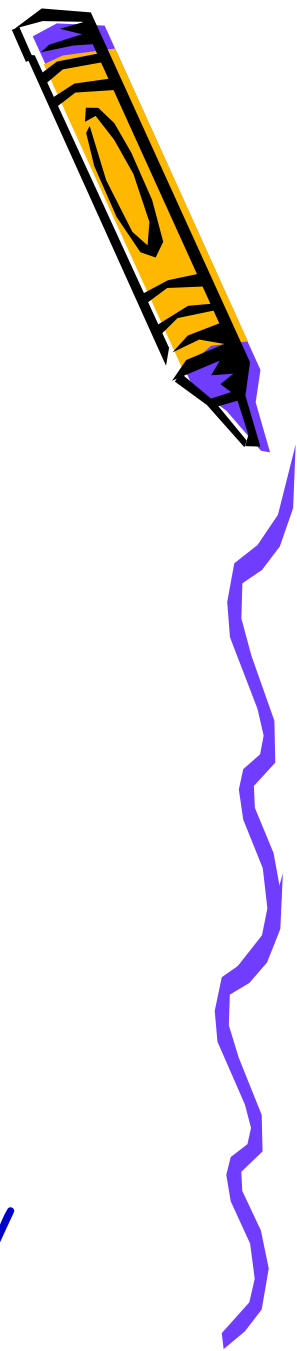
CPU Time (cont.)



- How these Parameters Determined?
 - Instruction count
 - CPI
 - Clock cycle time



CPU Time (cont.)



- Instruction Count (IC)
 - Determined by program
 - ISA
 - Compiler
- CPI
 - Determined by uArch
 - The way processor is implemented
 - Code CPI also depends on program
- Clock Cycle Time
 - Determined by uArch and Technology



Easy Practice



- Assume:
 - A C program compiled on two computers
 - Computer A with a RISC ISA
 - Computer B with a CISC ISA
 - Q1: which one would have higher IC?
 - Q2: which one would have smaller CPI?
 - Q3: which one would have higher performance?
 - Q4: can we have CPI less than one?



CPI Classes



- Question:
 - Do all instructions have same CPI?
 - No

| | CPI for instruction classes | | |
|-----|-----------------------------|---|---|
| | A | B | C |
| CPI | 1 | 2 | 3 |

| Code Sequence | Instruction count | | |
|---------------|-------------------|---|---|
| | A | B | C |
| 1 | 2 | 1 | 2 |
| 2 | 4 | 1 | 1 |



CPI Classes (cont.)



- CPU Clock Cycles

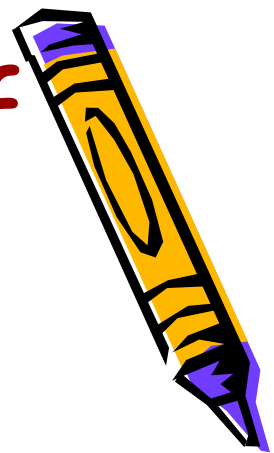
$$= CPI_1 * IC_1 + CPI_2 * IC_2 + ... + CPI_n * IC_n$$

| | CPI for instruction classes | | |
|-----|-----------------------------|---|---|
| | A | B | C |
| CPI | 1 | 2 | 3 |

| Code Sequence | Instruction count | | |
|---------------|-------------------|---|---|
| | A | B | C |
| 1 | 2 | 1 | 2 |
| 2 | 4 | 1 | 1 |



Performance Evaluation of Computers



- So Far
 - Learnt how to measure CPU Performance
- Consider Two CPUs
 - CPU1 runs faster than CPU2 on program A
 - CPU2 runs faster than CPU1 on program B
- Questions:
 - How we can decide which CPU is faster?
 - Which candidate programs should we choose to compare CPU1 and CPU2?



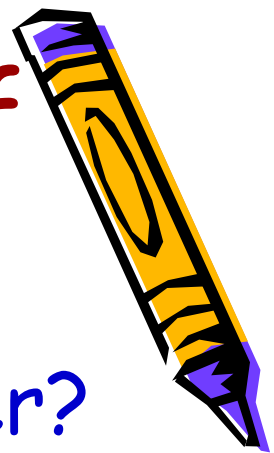
Performance Evaluation of Computers (cont.)



- Assume
 - A user typically runs programs A, B, and C in his computer (CPU1)
 - "Workload"
 - Set of programs A, B, and C
 - If CPU2 runs this workload faster
 - $\rightarrow \text{Performance}(\text{CPU2}) > \text{Performance}(\text{CPU1})$
 - In reality, we use a set of programs called **benchmarks** to compare performance of processors



Performance Evaluation of Computers (cont.)

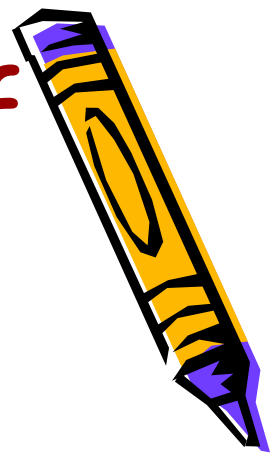


- Which CPU Runs a Workload Faster?
 - Depends on type of average we take over execution times
 - Simple arithmetic mean
 - Weighted arithmetic mean
 - Average over ratios
 - Geometric mean

| CPI | Computer A | Computer B |
|------------|------------|------------|
| Program 1 | 1 | 10 |
| Program 2 | 1000 | 100 |
| Total Time | 1001 | 110 |



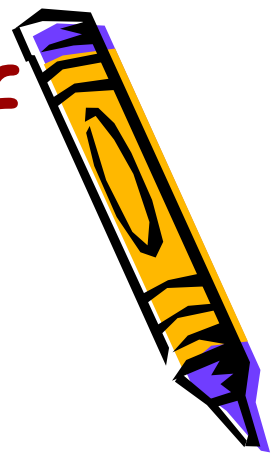
Performance Evaluation of Computers (cont.)



- Another Question:
 - Why not just running one simple program to compare performance of CPUs?
 - Agreeing on one simple program very hard
 - Designers can optimize processors towards fast running that simple program



Performance Evaluation of Computers (cont.)



- Which Programs to Choose?
 - Real programs such as MS-Word, Internet explorer, latex compilers
 - Synthetic benchmarks
 - Emulate frequency of different instructions in real programs
 - Standard benchmarks
 - Examine processor and memory hierarchy



Performance Evaluation of Computers (cont.)



- Which Programs to Choose?
 - Programs from different applications
 - SPEC CPU2000 benchmarks

| Integer Benchmarks (12) | | FP Benchmarks (14) | |
|-------------------------|-------------------------|--------------------|--|
| Name | Description | Name | Description |
| gzip | Compression | ammp | Computation chemistry |
| vpr | FPGA circuit P&R | swim | Shallow water model |
| gcc | C compiler | art | Image recognition using neural network |
| parser | Word processing program | galgel | Computational fluid dynamics |



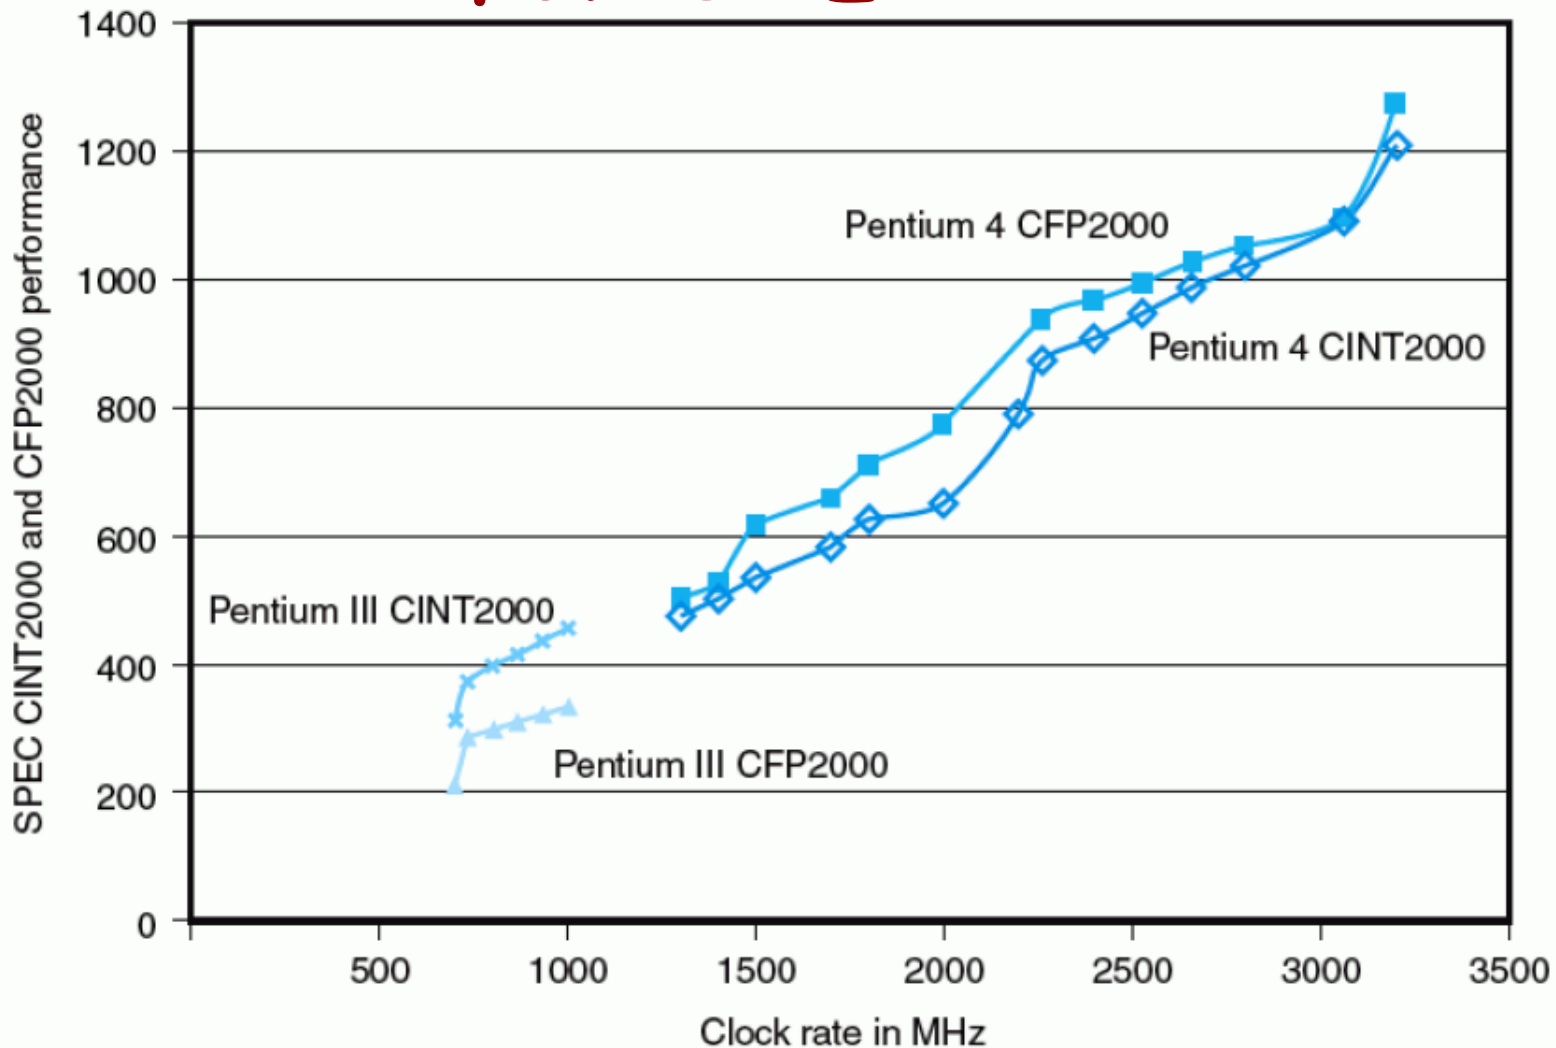
Easy Practice



- Question:
 - Can microprocessor designers optimize their architecture to optimize performance of SPEC2000 or other benchmark program?



Performance of PIII and P4 for SPEC2000



Experiment Setup



- How to Report Experiment Setup?
 - CPU frequency is NOT enough!
 - Need to Report both HW & SW config.
- Software
 - Operating system: WinXP prof. SP1
 - Compiler: Microsoft Visual Studio.NET
 - 7.0.xxx
 - File system type: NTFS
 - System state: default
 - Benchmark, Program, Input to program



Experiment Setup (cont.)



- Hardware
 - Hardware vender: Dell
 - Model number: Precision WorkStation 360
 - CPU: Intel Pentium 4 (800 MHz system bus)
 - CPU MHz: 3200
 - FPU: Integrated
 - Primary cache: 12KB (I), 8KB (D), both on-chip
 - Secondary cache: 512KB (I+D), on-chip
 - L3 cache: 2048KB (I+D), on-chip
 - Memory: 4x512MB ECC DDR400 SDRAM CL3
 - Disk subsystem: 1x80GB ATA/100 7200 RPM



CPU/Memory/IO Intensive Benchmarks



- SPEC Benchmarks
 - CPU/Memory intensive benchmarks
 - Some stress CPU
 - Some stress memory subsystem
- SPEC-Web
 - I/O intensive benchmarks
 - Mostly stress I/O subsystem
 - Disk subsystem, network connections, ...



SPECWeb99 Performance for Variety of Systems



| System | Processor | Number of disk drives | Number of CPUs | Number of networks | Clock rate (GHz) | Result |
|-----------|-------------------|-----------------------|----------------|--------------------|------------------|--------|
| 1550/1000 | Pentium III | 2 | 2 | 2 | 1 | 2765 |
| 1650 | Pentium III | 3 | 2 | 1 | 1.4 | 1810 |
| 2500 | Pentium III | 8 | 2 | 4 | 1.13 | 3435 |
| 2550 | Pentium III | 1 | 2 | 1 | 1.26 | 1454 |
| 2650 | Pentium 4 Xeon | 5 | 2 | 4 | 3.06 | 5698 |
| 4600 | Pentium 4 Xeon | 10 | 2 | 4 | 2.2 | 4615 |
| 6400/700 | Pentium III Xeon | 5 | 4 | 4 | 0.7 | 4200 |
| 6600 | Pentium 4 Xeon MP | 8 | 4 | 8 | 2 | 6700 |
| 8450/700 | Pentium III Xeon | 7 | 8 | 8 | 0.7 | 8001 |



Speedup




- $\text{Speedup} = \text{time}_{\text{original}} / \text{time}_{\text{improved}}$
- Example
 - $\text{time}_{\text{original}} = 100 \text{ s}$
 - $\text{time}_{\text{improved}} = 98 \text{ s}$
 - $\text{Speedup} = 100/98 = 1.02$



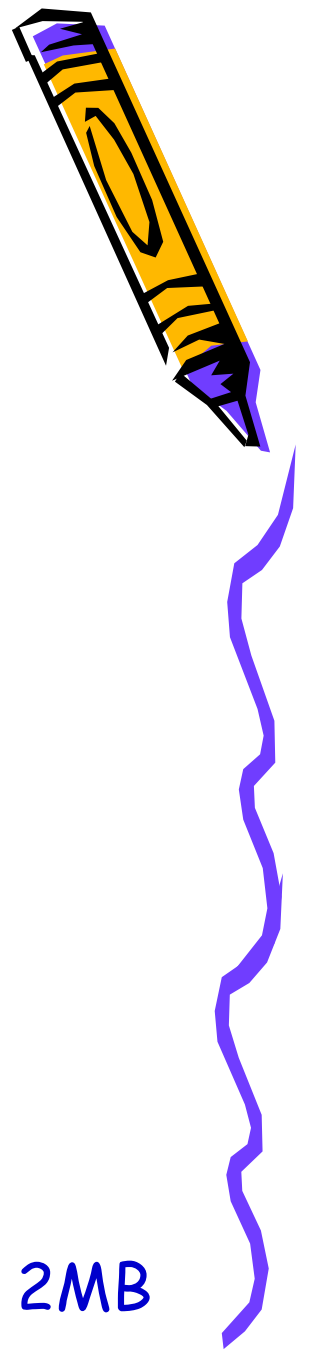
Amdahl's Law



- Amdahl's Law Says
 - Speedup limited to fraction improved
 - Obvious, but fundamental observation



Amdahl's Law (cont.)



- Obvious but Common Mistakes
 - CPU upgraded from 1.5Ghz to 3Ghz
 - But not seeing 100% improvement in performance
 - Combinational upgrade 1
 - CPU from 1.5Ghz to 3Ghz
 - Memory from 1GB-90nm to 2GB-45nm
 - Combinational upgrade 2
 - CPU from 1.5Ghz to 3Ghz
 - Memory from 1GB-90nm to 2GB-45nm
 - L1 from 16KB to 32KB & L2 from 1MB to 2MB



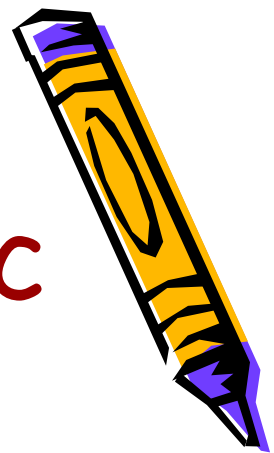
Speedup (cont.)

ExecTime_{new}

$$= \text{ExecTime}_{\text{old}} \times \{(1 - \text{Frac}_{\text{enhanced}}) + (\text{Frac}_{\text{enhanced}} / \text{Speedup}_{\text{enhanced}})\}$$

$$\text{Speedup}_{\text{overall}} = \text{ExecTime}_{\text{old}} / \text{ExecTime}_{\text{new}} = \frac{1}{(1 - \text{Frac}_{\text{enhanced}}) + (\text{Frac}_{\text{enhanced}} / \text{Speedup}_{\text{enhanced}})}$$





MIPS: Performance Metric

- MIPS
 - Million Instructions Per Second =
$$\text{Instruction Count} / (\text{Execution time} \times 10^6)$$
- Drawbacks?



MIPS:

Performance Metric (cont.)

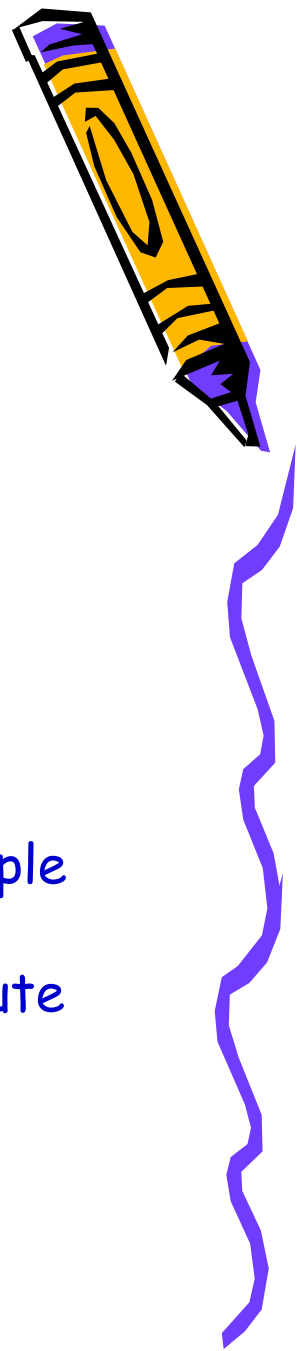


- Drawbacks
 - Not taking into account capabilities of instructions
 - Comparing two different ISAs not fair
 - Not realistic metric even on same CPU
 - Two programs: program A and program B
 - $MIPS(A) > MIPS(B) \rightarrow Perf.(A) > Perf.(B) ???$
 - Some optimization tech. add more code



MIPS:

Performance Metric (cont.)



- Example:
 - Machine A
 - Special instruction for performing square root
 - It takes 100 cycles to execute
 - Machine B
 - Doesn't have special instruction
 - must perform square root in software using simple instructions
 - e.g, Add, Mult, Shift each take 1 cycle to execute
 - Clock cycle = 1 μ s
 - Machine A: 1/100 MIPS = 0.01 MIPS
 - Machine B: 1 MIPS



MFLOPS: Performance Metric



- MFLOPS
= (FP ops/program) \times (program/time) $\times 10^{-6}$
- Popular in scientific computing
 - FP ops were previously much slower than regular instructions (i.e., off-chip, sequential execution)
- Not great for “predicting” performance
 - Ignores other instructions (e.g., load/store)
 - Not all FP ops have common format
 - Depends on how FP-intensive program is



Backup

